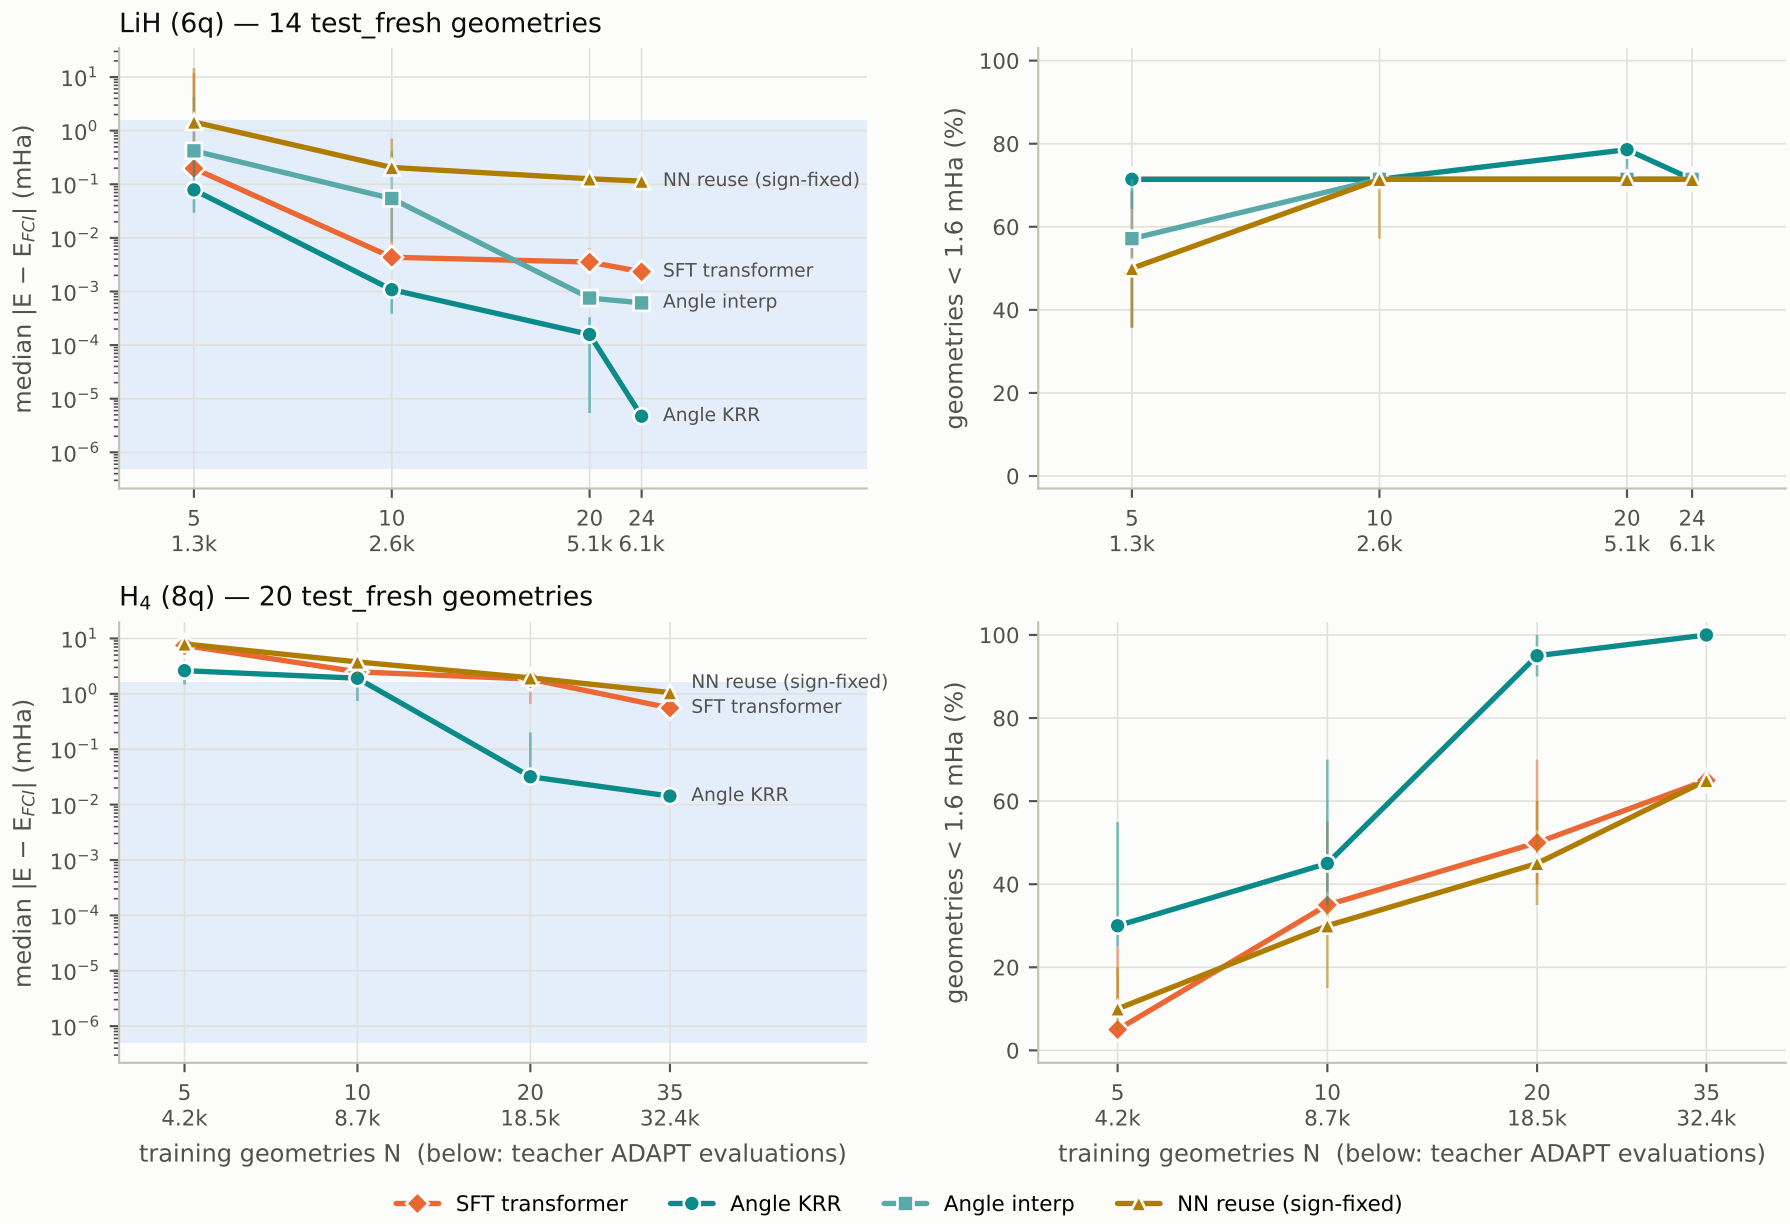

Nested random TRAIN subsets, 3 subset seeds per N (line = median over seeds, bar = seed range; the largest N is the full TRAIN set, one run). Teacher evaluations = ADAPT pool-gradient + energy evaluations summed over the N geometries. All methods one-shot, 0 energy evaluations at inference. Shaded: below 1.6 mHa.